CS:APP2e Web Aside ASM:OPT:
Machine Code Generated with Higher Levels of Optimization[*]

Randal E. Bryant
David R. O'Hallaron

July 21, 2009

Notice

1 Introduction

In the presentation of x86 machine code in CS:APP2e Chapter 3, we looked at machine code generated with level-one optimization (specified with the command-line option '-O1'.) In practice, most heavily used programs are compiled with higher levels of optimization. For example, all of the GNU libraries and packages are compiled with level-two optimization, specified with the command-line option '-O2'.

Recent versions of GCC employ an extensive set of optimizations at levels two and above. These optimizations can significantly improve program performance, but they make the mapping between source and machine code much more difficult to discern. This can make the programs more difficult to debug. Nonetheless, these higher level optimizations have now become standard, and so those who study programs at the machine level must become familiar with the possible optimizations they may encounter.

In this note, we describe some of the transformations performed by GCC to achieve higher levels of program optimization, and we explain the concepts behind them. Keep in mind, however, that GCC continues to evolve and become more sophisticated. We can anticipate that future versions of GCC will perform even more extensive transformations and optimizations than we cover here. The online documentation for GCC

1

[2] describes the many different optimizations that GCC performs (although without many details), and the different command-line arguments that control code generation.

It is important to recognize that there are other compilers for x86, particularly those supported by Intel and Microsoft. These compilers perform different optimization and transformations and have different code generators. When examining the machine code generated by an unfamiliar compiler, it is important to study the different code patterns the compiler generates. All of the examples we show here are compiled for IA32. Similar optimizations are made for x86-64.

2 Instruction Reorderings

(a) C code

```
1 int select(int x, int y, int i) {
2     int data[2] = { x, y };
3     if (i >= 0 && i < 2)
4         return data[i];
5     else
6         return 0;
7 }
```

(b) Generated assembly code, optimized -O1

```
    Function select, optimized -O1
    x at %ebp+8, y at %ebp+12, i at %ebp+16
 1 select:
    Setup code
 2   pushl   %ebp                    Save frame pointer                (S)
 3   movl    %esp, %ebp              Create new frame pointer          (S)
 4   subl    $16, %esp               Allocate 16 bytes on stack        (S)
    Body code
 5   movl    16(%ebp), %edx          Get i                             (B)
 6   movl    $0, %eax                Set result = 0                    (B)
 7   cmpl    $1, %edx                Compare i:1                       (B)
 8   ja      .L3                     if i < 0 or i > 1, goto done      (B)
 9   movl    8(%ebp), %eax           Get x                             (B)
10   movl    %eax, -8(%ebp)          Store in a[0]                     (B)
11   movl    12(%ebp), %eax          Get y                             (B)
12   movl    %eax, -4(%ebp)          Store in a[1]                     (B)
13   movl    -8(%ebp,%edx,4), %eax   Set result = a[i]                 (B)
14 .L3:                             done:
    Completion code
15   leave                          Restore stack and frame pointers  (C)
16   ret                            Return                            (C)
```

Figure 1: **Sample function and generated code with level-one optimization.** The code has separate sections for setup (S), body (B), and completion (C).

In the example functions for which we have examined the assembly code generated by GCC, there were several patterns that made it simpler to understand the code. Figure 1 shows an example of a simple C function and the code generated by GCC with level-one optimization. Two features are worth noting:

- The code for a function has distinct sections. First, there is a *setup* section, where the stack frame is created. These lines are labeled '(S)' on the right-hand side. Second, the *body* section performs the actual computations for the function, shown with lines labeled '(B).' Finally, the *completion* section deallocates the space for the stack frame and restores registers to their original values, with lines labeled '(C).'

- The function has a single return point, that is, there is only one `ret` instruction. Regardless of any branches in the code, all execution paths end with this instruction.

```
   Function select, optimized -O3
   x at %ebp+8, y at %ebp+12, i at %ebp+16
 1 select:
 2   pushl   %ebp                      Save frame pointer                   (S)
 3   xorl    %eax, %eax                Set result = 0                       (B)
 4   movl    %esp, %ebp                Create new frame pointer             (S)
 5   subl    $16, %esp                 Allocate 16 bytes on the stack       (S)
 6   movl    16(%ebp), %edx            Get i                                (B)
 7   cmpl    $1, %edx                  Compare i:1                          (B)
 8   jbe     .L6                       If i >= 0 and i <= 1, goto index     (B)
 9   leave                             Restore stack and frame pointers     (C)
10   ret                               Return                               (C)
11   .p2align 4,,7                     (Inserted to improve instruction alignment)
12   .p2align 3
13 .L6:                        index:
14   movl    8(%ebp), %eax             Get x                                (B)
15   movl    %eax, -8(%ebp)            Store in a[0]                        (B)
16   movl    12(%ebp), %eax            Get y                                (B)
17   movl    %eax, -4(%ebp)            Store in a[1]                        (B)
18   movl    -8(%ebp,%edx,4), %eax     Set result = a[i]                    (B)
19   leave                             Restore stack and frame pointers     (C)
20   ret                               Return                               (C)
```

Figure 2: **Function of Figure 1 and generated code with level-3 optimization.** The different sections are no longer distinct, and there are multiple return points.

Figure 2 shows the code generated for the same function, but with level-three optimization. (The same code is generated with optimization level 2.) We note that the overall code is basically the same, with several minor differences:

- The three sections of code are no longer distinct. We see the first instruction of the body (line 3) precedes some of the setup instructions.

- There are two return points, corresponding to the two branches of the conditional statement in the source code.

- There are directives in the code specifying *alignment constraints* (lines 11–12.) These directives constrain the following instructions to start at addresses that are multiples of some power of 2. The assembler satisfies this constraint by inserting bytes (with default value 0) into the instruction sequence. Since these bytes follow a `ret` instruction, they will never be encountered as instruction bytes by the executing program.

  The directive `.p2align` $n,,k$, indicates that the lower n bits of the address should be zero, and hence the address should be a multiple of 2^n, but only if this can be achieved by inserting at most k bytes. If it would require more than k bytes to achieve the desired alignment, then no bytes are inserted. The directive `.p2align` n directs it to insert as many bytes as are needed (at most $2^n - 1$) to achieve the desired alignment.

  As this example illustrates, these alignment directives are inserted to line up the starting address of a jump or call target. This can improve the speed at which the processor fetches instructions from memory.

  Practice Problem 1:

  Suppose the address following the `ret` instruction on line 10 of Figure 2 is of the form $16a + b$, where $0 \leq b < 16$.

  A. Suppose just directive `.p2align 4,,7` is present. How many bytes will be inserted, as a function of b?

  B. Suppose just directive `.p2align 3` is present. How many bytes will be inserted, as a function of b?

  C. Suppose both directives `.p2align 4,,7` and `.p2align 3` are present. How many bytes will be inserted, as a function of b?

These differences between level-one and higher levels of optimization are fairly minor. The actual instructions that are executed will be the same in either case, but in a slightly different order, and with two return points. In generating this code, GCC is attempting to optimize *instruction scheduling*, the order in which instructions are executed. The general strategy is to order instructions so that there is some gap between the instruction that generates a value and any instruction that uses that value. For this example, the modifications seen in the more highly optimized code would not lead to any real performance improvement. In fact, instruction scheduling has only limited benefit for programs running on modern processors, because they employ *out-of-order* execution, meaning that they dynamically reorder the instructions to maximize performance. These processors and their performance characteristics are described in CS:APP2e Chapter 5.

3 Program Example

For the remaining part of this note, we will use functions that operate on singly-linked lists as our program examples. The declaration of the list data structure and some basic functions operating on lists are shown in Figure 3. Linked-list code provides good examples for program optimization, especially since many list operations are naturally written as recursive functions.

(a) List data type declaration

```
1 /* Linked list element */
2 typedef struct ELE {
3     int val;
4     struct ELE *next;
5 } list_ele, *list_ptr;
```

(b) List access functions

```
1 int val(list_ptr ls) {
2     if (ls == NULL)
3         return 0;
4     else
5         return ls->val;
6 }
7
8 list_ptr next(list_ptr ls) {
9     if (ls == NULL)
10        return NULL;
11    else
12        return ls->next;
13 }
14
15 int is_null(list_ptr ls) {
16    return ls == NULL;
17 }
```

Figure 3: **Singly-linked list code examples.** We will use functions on lists to demonstrate optimizations made by GCC.

Figure 4 shows the C code for a function that computes the sum of the elements in a list, written using iteration. The assembly code generated is also shown. This code was generated with level-three optimization, but similar code occurs for levels one and two. Note carefully the loop, consisting of four instructions (lines 10–14.)

4   Inline Substitution

Inline substitution is a very basic and effective technique for eliminating the overhead due to function calls and to enable optimizations based on special characteristics of the context in which a function is called. It involves replacing the call to a function with the code that implements that function.

Figure 5 gives a demonstration of inline substitution. First, we show a version of the list summation function (a) in which we use the functions is_null, val, and next (shown in Figure 3), rather than directly accessing the list elements. These functions are typical of the *method* functions found in object-oriented languages such as C++ and Java. They involve very short segments of code, and they often include error checking, such as checking for null pointers.

We show the result of substituting the code for the functions is_null, val, and next into the summation code in part (b). By compiling this version, we would eliminate the overhead of calling the list functions, including the effort required to allocate and deallocate their stack frames. More significantly, the compiler is able to optimize away some of the redundant code shown in part (b). In particular, based on the fact that the loop will exit when variable ls equals NULL, the compiler can determine that ls will be non-null inside the loop. That means the two tests ls == NULL, and the statements for the "then" cases can be optimized away. The net effect of these substitutions plus optimizations is that the code generated for this

(a) C code

```
1 int sum_list_iter(list_ptr ls) {
2     int sum = 0;
3     while (ls != NULL) {
4         sum += ls->val;
5         ls = ls->next;
6     }
7     return sum;
8 }
```

(b) Generated assembly code

```
   Code generated for iterative version of list sum
   with optimization level -O3
   ls at %ebp+8
1 sum_list_iter:
2   pushl   %ebp
3   xorl    %eax, %eax        Set sum = 0
4   movl    %esp, %ebp
5   movl    8(%ebp), %edx     Get ls
6   testl   %edx, %edx
7   je      .L62              if ls == NULL, goto done
8   .p2align 4,,7             (Inserted to improve instruction alignment)
9   .p2align 3
   Within loop: ls in %edx, sum in %eax
10 .L65:                      loop:
11   addl    (%edx), %eax     Add ls->val to sum
12   movl    4(%edx), %edx    Set ls = ls->next
13   testl   %edx, %edx
14   jne     .L65             If ls == NULL, goto loop
15 .L62:                      done:
16   popl    %ebp
17   ret
```

Figure 4: **Iterative summation of list.** Similar code is generated for all optimization levels.

(a) C code

```
1  int sum_list_iter_abs(list_ptr ls) {
2      int sum = 0;
3      while (!is_null(ls)) {
4          sum += val(ls);
5          ls = next(ls);
6      }
7      return sum;
8  }
```

(b) C code with functions expanded by inline substitution

```
1  /* Result of inline expansion in function sum_list_iter_abs */
2  int sum_list_iter_expand(list_ptr ls) {
3      int sum = 0;
4      /* Expansion of function is_null */
5      while (!(ls==NULL)) {
6          {
7              /* Expansion of function val */
8              int val;
9              if (ls == NULL) /* Optimized away */
10                 val = 0;
11             else
12                 val = ls->val;
13             sum += val;
14         }
15         {
16             /* Expansion of function next */
17             list_ptr next;
18             if (ls == NULL) /* Optimized away */
19                 next = NULL;
20             else
21                 next = ls->next;
22             ls = next;
23         }
24     }
25     return sum;
26 }
```

Figure 5: **Demonstration of inline substitution.** The compiler is then able to generate code identical to that for sum_list_iter (Figure 4.)

more abstract version of list summation is identical to that generated for the less abstract version (Figure 4.) Writing code in a more abstract style, such as is shown in Figure 5, need not incur any performance penalty.

(a) C code

```
1 int test_select() {
2     return select(5, 6, 1);
3 }
```

(b) Generated assembly code, optimized -O1

```
   Function test_select, optimized -O1
 1 test_select:
 2   pushl   %ebp
 3   movl    %esp, %ebp
 4   subl    $12, %esp                Allocate 12 bytes on stack
 5   movl    $1, 8(%esp)              Set 1 as 3rd argument
 6   movl    $6, 4(%esp)              Set 6 as 2nd argument
 7   movl    $5, (%esp)               Set 5 as 1st argument
 8   call    select                   Call select(5,6,1)
 9   leave
10   ret
```

(c) Generated assembly code, optimized -O3

```
   Function select, optimized -O3
   x at %ebp+8, y at %ebp+12, i at %ebp+16
 1 test_select:
 2   pushl   %ebp
 3   movl    $6, %eax                 Set result = 6
 4   movl    %esp, %ebp
 5   popl    %ebp
 6   ret
```

Figure 6: **Example of inline substitution followed by optimization.** GCC determines that the function will always return 6.

Figure 6 shows an interesting example of the performance benefits to be gained by inline substitution. In (a), we show an example function that calls the function select, shown in Figure 1. When run with optimization level one (b), GCC does not perform inline substitution, and hence it has no choice but to set up a call to select. By performing inline substitution as part of level-three optimization (c), the compiler is able to detect that the behavior of this particular call to test_select is highly predictable, and in fact will always return 6. So, the compiler simply generates code that sets register %eax to 6.

Drawbacks and Limitations of Inline Substitution

Our examples show the performance advantages of inline substitution. Eliminating the overhead of function calls, including allocating and deallocating a stack frame can be significant. More importantly, inline substitution enables a context-based optimization of the function code.

On the other hand, there are several important drawbacks and limitations of inline substitution:

The code is harder to debug. Inline substitution eliminates some of the call and return behavior expressed in the source code. This normally doesn't matter, because the computed results will be identical. However, if we try to monitor the executable program with a debugger such as GDB, we could find some unexpected results. For example, if we set a breakpoint for the function `is_null` in the list code, we might be surprised that the sum function completes without ever hitting the breakpoint. A general rule of thumb is to set the optimization level lower when debugging a program, and increase it only when generating production code.

The code size can grow. Each inline substitution can cause a replication of an entire function's worth of code. In the worst case, it can even cause a program to blow up to a size that is exponential in the size of the original source program. Compilers have complex heuristic rules to decide whether or not to perform inline substitution for a given function. These rules typically err on the side of caution, in order to keep the generated code within a tight bound of what would be generated without inline substitution.

It may not be possible. Inline substitution requires that the compiler have access to the source code of a function at compile time. This was possible for our list programs, because they were all within a single file. Ordinarily, however, library functions are precompiled, and large programs are divided into multiple files that are compiled separately. The compiler might have access to the function prototypes, via a '.h' file, but not to the source code. Putting everything in one file runs counter to the goal of making programs as modular as possible.

5   Recursion Removal

Many programmers feel that expressing a computation as a recursive function can provide a clearer depiction of the desired behavior. It is a natural expression of a "divide and conquer" strategy, where we divide a problem into pieces, solve them separately (by recursive calls), and then combine the results.

From a performance perspective, however, recursive functions have two drawbacks. First, they tend to run slower. For example, we measured different factorial functions on an Intel Core i7 processor and found the iterative version runs around twice the speed of the recursive one. Second, they require more space. We have seen, for example, that our recursive factorial program allocates 16 bytes on the stack for each call, and these allocations accumulate until the termination condition is reached. Thus, computing the factorial of a value n recursively will require around $16n$ bytes of stack space. By contrast, the iterative version of factorial requires 8 bytes of stack space regardless of the value of n.

Some computations fundamentally require more than a constant amount of storage, and hence some recursive functions cannot be expressed as iterative computations, without adding an additional data structure such as a stack. For the class of *linear* recursions, however, where any invocation of a function contains at most one recursive call, we can often transform recursion into iteration. When this can be done automatically by the compiler, we gain the advantage of allowing the programmer to express a function in a clear manner and having the compiler transform it into a form that requires less time and space.

5.1 Tail Recursion

The most straightforward form of recursion is referred to as *tail recursion*. A call by a function f to a function g is labeled as a *tail call* when the result of this call to g is returned directly as the result of function f. A *tail-recursive* function is one for which all recursive calls are tail calls.

As an example, the following function computes the sum of a linked list using tail recursion.

```
1  int sum_list_tail(list_ptr ls, int sofar) {
2      if (ls == NULL)
3          return sofar;
4      else
5          return sum_list_tail(ls->next, sofar + ls->val);
6  }
7
8  int sum_list_tail_call(list_ptr ls) {
9      return sum_list_tail(ls, 0);
10 }
```

At a top level, the user would invoke the function sum_list_tail_call to compute the sum of a list, which in turn calls the recursive function sum_list_tail. We see that sum_list_tail is tail recursive—it contains only a single call to itself, and the result of this call is returned as the function result. Function sum_list_tail employs a strategy commonly seen in tail-recursive functions, where an additional argument serves as an "accumulator," here named sofar, to keep track of the sum of all of the list elements encountered up to this point. This variable is initialized to zero, and the function keeps adding values to it as it traverses the list. Once we reach the end of the list, the function can simply return sofar as the list sum.

A tail-recursive function can be automatically transformed into an iterative one. We describe the process here by first giving a more abstract representation of a tail-recursive function. The following code shows the general structure of a two-argument, tail-recursive function. (The same idea can be used for a function of any number of arguments.)

```
tail(a₁, a₂) {
    if (Cond(a₁, a₂)) {
        return Result(a₁, a₂);
    else {
        na₁ = Nval₁(a₁, a₂);
        na₂ = Nval₂(a₁, a₂);
        return tail(na₁, na₂);
    }
}
```

The general from has arguments a_1 and a_2. It checks whether it has reached a terminal case based on some condition $Cond$, applied to the arguments. The terminal condition returns a value $Result$, which can depend on the arguments. The recursive call involves computing new values for the arguments, based on computations $Nval_1$ and $Nval_2$, and then making a tail call.

For our list sum function, the correspondence between the general form and the actual code is based on the following mappings:

Abstract value	Program value
a_1	`ls`
a_2	`sofar`
$Cond(a_1, a_2)$	`ls == NULL`
$Result(a_1, a_2)$	`sofar`
$Nval_1(a_1, a_2)$	`ls->next`
$Nval_2(a_1, a_2)$	`sofar + ls->val`

We can transform the general, tail-recursive function into a version that uses iteration as follows:

```
itail(a₁, a₂) {
    while (!Cond(a₁, a₂)) {
        na₁ = Nval₁(a₁, a₂);
        na₂ = Nval₂(a₁, a₂);
        a₁ = na₁;
        a₂ = na₂;
    }
    return Result(a₁, a₂);
}
```

In this code, we keep updating the variables a_1 and a_2 in the same manner as we would via repeated calls to the recursive function. (We show this as first computing values na_1 and na_2 and then assigning these to a_1 and a_2 to avoid any inconsistency of having the computation of $Nval_2$ use the updated version of a_1.) Once we reach the terminating condition, we can simply return the value computed by the base case.

If we apply this transformation to the tail-recursive list summation example, we get the following code:

```
1  int sum_list_itail(list_ptr ls, int sofar) {
2      while (!(ls == NULL)) {
3          list_ptr nls    = ls->next;
4          int      nsofar = sofar + ls->val;
5          ls    = nls;
6          sofar = nsofar;
7      }
8      return sofar;
```

```
9  }
```

We can see that the inner loop of this function is essentially the same as for the iterative sum function (Figure 4). In fact, using a combination of inline substitution and tail recursion removal, GCC generates the same code for top-level function `sum_list_tail_call`, as it does for the iterative sum function `sum_list_iter`.

> **Practice Problem 2**:
>
> The following is a recursive factorial function. We assume that argument `x` is greater than or equal to zero.
>
> ```
> 1  int fact_recur(int x) {
> 2      if (x == 0)
> 3          return 1;
> 4      else
> 5          return x * fact_recur(x-1);
> 6  }
> ```
>
> A. Write a version of the function based on tail recursion. Show how this function would be called to compute the factorial of `x`.
>
> B. Show a mapping between your function and the general form of a tail-recursive function.
>
> C. Show how your tail-recursive function would be transformed into an iterative one.

Tail recursion elimination is a standard optimization used by many compilers [4]. Its only drawback is when we then try to monitor code execution via a debugger such as GDB. If we set a breakpoint for the function, we might be surprised to find the breakpoint only gets triggered once. In fact, if the compiler uses a combination of inline substitution and tail-call elimination, we would find that calling the function `sum_list_tail_call` does not trigger the breakpoint for `sum_list_tail` at all.

5.2   More General Forms of Recursion

Many recursive functions do not use tail recursion. In fact, we can see for our list summation example that tail recursion was not the most natural way to express the desired computation. Instead, a more traditional way of expressing list summation is via the recursive function `sum_list_rec` shown in Figure 7(a). As shown in part (b) of the figure, GCC is able to transform this recursive function into an iterative computation. We see that the generated code is almost identical to that for the iterative summation (Figure 4). The only difference is that the inner loop requires 5 instructions to do the same computations as does the earlier version do with 4 instructions.

Being able to transform recursion into iteration is a more advanced form of optimization, not covered in most books on compiler optimization. Our presentation is based on a research paper [3].

Suppose the argument list consists of elements $l_1, l_2, \ldots, l_{n-1}, l_n$, where l_1 is the first element of the list. Then our iterative summation function computes the sum as:

$$[\cdots [[0 + l_1] + l_2] + \cdots + l_{n-1}] + l_n \tag{1}$$

(a) C code

```
1 int sum_list_rec(list_ptr ls) {
2     if (ls == NULL)
3         return 0;
4     else
5         return ls->val + sum_list_rec(ls->next);
6 }
```

(b) Generated assembly code, optimized -O3

```
    Code generated for recursive version of list sum
    with optimization level -O3
    ls at %ebp+8
 1 sum_list_rec:
 2   pushl    %ebp
 3   xorl     %ecx, %ecx          Set sum = 0
 4   movl     %esp, %ebp
 5   movl     8(%ebp), %edx       Get ls
 6   testl    %edx, %edx
 7   je       .L14                If ls == NULL, goto done
 8   .p2align 4,,7               (Inserted to improve instruction alignment)
 9   .p2align 3
    Within loop: ls in %edx, sum in %ecx
10 .L17:                          loop:
11   movl     (%edx), %eax        Get val = ls->val
12   movl     4(%edx), %edx       Set ls = ls->next
13   addl     %eax, %ecx          Add val to sum
14   testl    %edx, %edx
15   jne      .L17                If ls != NULL, goto loop
16 .L14:                          done:
17   movl     %ecx, %eax          Set result = sum
18   popl     %ebp                Restore frame pointer
19   ret                          Return
```

Figure 7: **Recursive list summation function and generated assembly code.** The compiler is able to transform the recursion into iteration, even though it is not tail recursive.

(We use brackets '[' and ']' in our expression to denote the grouping of terms. This serves to highlight the difference between grouping and the use of parentheses to denote function application.) We can see by this equation that the computation starts with a sum value of zero and adds successive list elements from l_1 to l_n.

Our recursive summation function computes the sum as:

$$l_1 + [l_2 + \cdots + [l_{n-1} + [l_n + 0]] \cdots] \tag{2}$$

That is, the sequence of recursive calls steps through the list until it reaches the empty list, at which point the function returns 0. Then the returns from the calls sum the list elements in reverse order from l_n to l_1. The two expressions yield identical results, because integer addition is both commutative and associative. By way of comparison, they might produce different results if we were using floating-point arithmetic, since it is not associative. Our goal is to find a general way to convert the first form of computation into the second, exploiting the mathematical properties of the operation used to generate the final results.

We can express the recursion seen in function `list_sum_rec` in a general form by the function `recur`:

```
recur(a) {
    if (a == a₀) {
        /* Base case */
        return b₀;
    else {
        /* Recursive call */
        return V(a) op recur(N(a));
}
```

This general form has a single argument a. The base case is reached when the argument matches a base value, a_0, in which case the function returns some value b_0. The recursive case involves computing a new argument $N(a)$, and making the recursive call. The result of this call is then combined with value $V(a)$ via a *combining operation op*.

For our list sum function, the correspondence between the general form and the actual code is based on the following mappings:

Abstract value	Program value
a	`ls`
a_0	`NULL`
b_0	`0`
$V(a)$	`ls->val`
$N(a)$	`ls->next`
op	`+`

Assume that we start the computation with argument a, and that $N^k(a) = a_0$, where the notation N^k means

k applications of function N. Then we can write the computation performed by the recursive function as:

$$V(a) \; op \; [V(N(a)) \; op \; \cdots \; op \; [V(N^{k-2}(a)) \; op \; [V(N^{k-1}(a)) \; op \; b_0]] \cdots] \qquad (3)$$

For the case where combining operation op is associative and commutative, as is integer addition, we can rearrange the terms of this expression into a form similar to our iterative version of list summation:

$$[\cdots [[b_0 \; op \; V(a)] \; op \; V(N(a))] \; op \; \cdots \; op \; V(N^{k-2}(a))] \; op \; V(N^{k-1}(a)) \qquad (4)$$

This expression can be converted to an iterative computation, much as we did for summing the elements of a list, introducing a variable r to hold the accumulated result. We label this the "iter1" transformation.

```
iter1(a) {
    r = b₀;
    while (a != a₀) {
        r = r op V(a);
        a = N(a);
    }
    return r;
}
```

We can see that when we apply this transformation to recursive list summation, we get the C code shown for `sum_list_iter` (Figure 4), where `sum` serves as the accumulating variable r. As we have seen, GCC generates nearly the same code for the functions `sum_list_iter` (Figure 4) and `sum_list_rec` Figure 7), an indication that it applies the `iter1`.

> **Practice Problem 3**:
>
> Let us apply the `iter1` transformation to convert the recursive factorial function, shown in Problem 2, into an iterative one.
>
> A. Show the mapping between the factorial code and the general recursive form of function `recur`. Is operation op both associative and commutative?
> B. Show the code you would get by applying the `iter1` transformation.

In our experience, GCC applies the `iter1` transformation when applicable. The main limitation is that it must be able to determine that the combining operation is commutative and associative. It applies the transformation for integer addition (as in list summation) and multiplication (as in factorial), but not for floating-point operations. In addition, it will not apply the transformation for more complex operations, even if they are commutative and associative. For example, consider code to find the maximum element in a list:

```
1 int max(int x, int y) {
2     return (x > y) ? x : y;
3 }
4
5 int max_list_rec(list_ptr ls) {
6     if (ls == NULL)
7         return INT_MIN;
8     else
9         return max(ls->val,
10                     max_list_rec(ls->next));
11 }
```

The code for `max_list_rec` matches the pattern for `recur`, with `max` serving as the combining opera-
tion. This operation is both commutative and associative, but GCC has no way of determining this.

Practice Problem 4:

Consider the "list difference" computation, where we replace the addition operation of the list summation
function with subtraction:

```
1 int diff_list_rec(list_ptr ls) {
2     if (ls == NULL)
3         return 0;
4     else
5         return ls->val - diff_list_rec(ls->next);
6 }
```

Unfortunately, subtraction is neither commutative nor associative.

A. Show how you can transform this function into one that uses addition as the combining operation,
   possibly introducing a second function argument. How would this be called at the top level?

B. Show the iterative form you would get by applying a transformation similar to `iter1`, but com-
   bining the top-level call with the two-argument recursive function.

Practice Problem 5:

Consider the following code, declaring a data type for the nodes of a binary tree, and a recursive function
to sum the values of all the nodes in a tree:

```
1 /* Tree data structure */
2 typedef struct NODE {
3     int val;
4     struct NODE *left;
5     struct NODE *right;
6 } tree_ele, *tree_ptr;
7
8 /* Sum values in tree, recursively */
9 int sum_tree_rec(tree_ptr tp) {
```

```
10      if (tp == NULL)
11          return 0;
12      else {
13          return tp->val +
14              sum_tree_rec(tp->left) +
15              sum_tree_rec(tp->right);
16      }
17 }
```

Function `sum_tree_rec` uses *nonlinear* recursion. In fact, this is an example of a computation that fundamentally requires a stack (such as is provided to support recursion) or some other data structure.

Even though it is not possible to eliminate both recursive calls, we can transform this function into one that computes the sum of the left subtree recursively and computes the sum of the right subtree iteratively.

A. Show a mapping between the code for `sum_tree_rec` and the general form `recur` that will allow the recursive call to the right subtree to be a candidate for the `iter1` transformation.

B. Show the C code you would get by applying the `iter1` transformation.

5.3   Partial Recursion Expansion

When GCC is unable to remove recursion altogether, it can still attempt to reduce some of the overhead of recursive function calls. It does this by partially expanding the recursion, generating code that steps through the first k recursive calls explicitly, for some value of k. This process can be viewed as a form of inline substitution—it replaces a recursive call with the body of the function a total of 8 times. This transformation is only applied for optimization levels 3 and higher, since it significantly increases the amount of code generated.

Let us examine how partial expansion can be applied to the list difference code shown in Problem 4. GCC is unable to transform this recursion into iteration, since subtraction is not associative. Figures 8–10 show the code GCC generates code for optimization level 3. This code steps through the first nine levels of recursion, storing the list elements either in registers or within the function's stack frame. After nine steps, it then performs a recursive call to sum the rest of the list. This recursive call, of course, will step through the next nine elements of the list before making a second recursive call, and so on, such that the total number of function calls is reduced by a factor of around nine.

Of course, the list may have fewer than eight elements, and so the code has branches after each step to bypass any remaining steps once it hits a null pointer. We show these branches as having destinations with labels of the form `len_i`, meaning the branch will be taken when the argument list is of length i. Observe how the explicit steps complete by stepping through the list elements in reverse order.

In our measurements of this transformation, we found that it typically speeds up the computation by a factor of around 1.2. It is not as effective as eliminating recursion altogether, but it does have some benefit.

Practice Problem 6:

Based on the code shown in Figures 8–10, show how to expand the general form of recursion given as function `recur` for $k = 3$ steps. Your code should not contain any `goto` statements. You should not

```
   Code generated for recursive version of list difference
   with optimization level -O3
   ls at %ebp+8
 1 diff_list_rec:
 2   pushl   %ebp
 3   xorl    %eax, %eax         Set diff to 0
 4   movl    %esp, %ebp
 5   subl    $40, %esp
   ls in register %edx, diff in register %eax
 6   movl    8(%ebp), %edx
 7   movl    %ebx, -12(%ebp)
 8   movl    %esi, -8(%ebp)
 9   movl    %edi, -4(%ebp)
10   testl   %edx, %edx
11   je      .L133              If ls==NULL, goto len_0
   Step through up to 8 elements of list, storing  values in registers and on stack
12   movl    (%edx), %esi       Set %esi to l_1
13   xorl    %eax, %eax         Set diff to 0
14   movl    4(%edx), %edx      ls = ls->next
15   testl   %edx, %edx
16   je      .L135              If ls==NULL, goto len_1
   ls in register %eax, diff in register %edx
17   movl    4(%edx), %eax      ls = ls->next
18   movl    (%edx), %edi       Set %edi to l_2
19   xorl    %edx, %edx         Set diff to 0
20   testl   %eax, %eax
21   je      .L137              IF ls==NULL, goto len_2
22   movl    (%eax), %edx
23   movl    4(%eax), %eax      ls = ls->next
24   movl    %edx, -36(%ebp)    Store l_3 at %ebp-36
25   xorl    %edx, %edx
26   testl   %eax, %eax
27   je      .L139              If ls==NULL, goto len_3
```

Figure 8: **Compilation of diff_list_rec (Problem 4), Part I.** This code demonstrates the expansion of recursion performed for optimization levels 3 and higher.

```
28   movl    (%eax), %edx
29   movl    4(%eax), %eax      ls = ls->next
30   movl    %edx, -32(%ebp)    Store l_4 at %ebp-32
31   xorl    %edx, %edx         Set diff to 0
32   testl   %eax, %eax
33   je      .L141              If ls==NULL, goto len_4
34   movl    (%eax), %edx
35   movl    4(%eax), %eax      ls = ls->next
36   movl    %edx, -28(%ebp)    Store l_5 at %ebp-28
37   xorl    %edx, %edx         Set diff to 0
38   testl   %eax, %eax
39   je      .L143              If ls==NULL, goto len_5
40   movl    (%eax), %edx
41   movl    4(%eax), %eax      ls = ls->next
42   movl    %edx, -24(%ebp)    Store l_6 at %ebp-24
43   xorl    %edx, %edx         Set diff to 0
44   testl   %eax, %eax
45   je      .L145              If ls==NULL, goto len_6
46   movl    (%eax), %edx
47   movl    4(%eax), %eax      ls = ls->next
48   movl    %edx, -20(%ebp)    Store l_7 at %ebp-20
49   xorl    %edx, %edx         Set diff to 0
50   testl   %eax, %eax
51   je      .L147              If ls==NULL, goto len_7
52   movl    (%eax), %edx
53   movl    4(%eax), %eax      ls = ls->next
54   movl    %edx, -16(%ebp)    Store l_8 at %ebp-16
55   xorl    %edx, %edx         Set diff to 0
56   testl   %eax, %eax
57   je      .L149              If ls==NULL, goto len_8
   Reach here only if list has >= 9 elements
58   movl    (%eax), %ebx       Set %ebx to l_9
59   movl    4(%eax), %eax      ls = ls->next
60   movl    %eax, (%esp)
   Make recursive call to compute difference for rest of list
61   call    diff_list_rec      diff = diff_list_rec(ls)
```

Figure 9: **Compilation of `diff_list_rec` (Problem 4), Part II.**

```
       Accumulate values for first nine elements of list, in reverse order
62   movl    %ebx, %edx        Retrieve l_9
63   subl    %eax, %edx        diff = l_9 - diff
64 .L149:                      len_8:
65   movl    -16(%ebp), %eax   Retrieve l_8
66   subl    %edx, %eax
67   movl    %eax, %edx        diff = l_8 - diff
68 .L147:                      len_7:
69   movl    -20(%ebp), %eax   Retrieve l_7
70   subl    %edx, %eax
71   movl    %eax, %edx        diff = l_7 - diff
72 .L145:                      len_6:
73   movl    -24(%ebp), %eax   Retrieve l_6
74   subl    %edx, %eax
75   movl    %eax, %edx        diff = l_6 - diff
76 .L143:                      len_5:
77   movl    -28(%ebp), %eax   Retrieve l_5
78   subl    %edx, %eax
79   movl    %eax, %edx        diff = l_5 - diff
80 .L141:                      len_4:
81   movl    -32(%ebp), %eax   Retrieve l_4
82   subl    %edx, %eax
83   movl    %eax, %edx        diff = l_4 - diff
84 .L139:                      len_3:
85   movl    -36(%ebp), %eax   Retrieve l_3
86   subl    %edx, %eax
87   movl    %eax, %edx        diff = l_3 - diff
88 .L137:                      len_2:
89   movl    %edi, %eax        Retrieve l_2
90   subl    %edx, %eax        diff = l_2 - diff
91 .L135:                      len_1:
92   subl    %eax, %esi        Retrieve l_1
93   movl    %esi, %eax        diff = l_1 - diff
94 .L133:                      len_0:
95   movl    -12(%ebp), %ebx
96   movl    -8(%ebp), %esi
97   movl    -4(%ebp), %edi
98   movl    %ebp, %esp
99   popl    %ebp
100  ret                       Return diff
```

Figure 10: **Compilation of `diff_list_rec` (Problem 4), Part III.**

assume any properties of the combining operation op; for example, it may be neither commutative nor associative.

5.4 *Other Recursion Eliminations: Beyond GCC

Although not currently implemented by GCC, it is interesting to consider other ways of converting functions having the recursive structure of `recur` into iterative ones. We explore these via a series of exercises.

Practice Problem 7:

Consider the case where the combining operation op is associative, but not necessarily commutative. (Examples of associative, but non-commutative operations include matrix multiplication and string concatenation.) By being a little more careful in how we order the computation, we can devise a scheme that will transform computations of the form shown in function `recurse` into iterative ones.

In devising the `iter1` transformation, we started with the expression (Equation 3) describing the computation performed by the recursive function `recurse`:

$$V(a) \ op \ [V(N(a)) \ op \cdots op \ [V(N^{k-2}(a)) \ op \ [V(N^{k-1}(a)) \ op \ b_0]] \cdots]$$

A. For an associative operation op, we can shift the parentheses in this expression, but we cannot change the positions of any of the elements. Show how we could reparenthesize the expression in a way that makes it suitable for an iterative computation.

B. Express the transformation to an iterative computation with a function `iter2` that is similar to `iter1`, except that it does not rely on commutativity. Be careful of the case where the argument a equals the base value a_0.

Practice Problem 8:

Show the function you would get by applying the `iter2` transformation you derived for Problem 7 to the recursive factorial function of Problem 2

Practice Problem 9:

Consider the case where the combining operation op is not associative. For example, subtraction is neither associative nor commutative, and neither floating-point addition nor multiplication are associative. Suppose on the other hand, that the function N we use to enumerate successive elements is invertible. That is, there is some function P such that for any value x, we have $P(N(x)) = x$. (This property does not apply to list summation, since there is no inverse for the expression `ls->next`.)

Consider the expression describing the computation performed by the recursive function `recurse` (Equation 3):

$$V(a) \ op \ [V(N(a)) \ op \cdots op \ [V(N^{k-2}(a)) \ op \ [V(N^{k-1}(a)) \ op \ b_0]] \cdots]$$

A. Rewrite this expression in terms of successive applications of function P, rather than N.

B. Convert your expression into an iterative function `iter3`. Be careful of the case where argument a equals the base value a_0.

Practice Problem 10:

Show how the `iter3`

transformation you derived in Problem 9 can be applied to the recursive factorial function of Problem 2.

    A. What is the inverse operation?

    B. What C code do you get by applying the `iter3` transformation?

6   Sibling Call Optimization

We saw that tail recursion can generally be replaced by iteration. A related form of optimization concerns *sibling calls*—tail calls to other functions. Suppose that function `f` makes a tail call to function `g`, as illustrated by the following code:

```
int f(int x)
{
    int y = P(x); /* Some computation */
    return g(y);  /* Tail call */
}
```

Once function `f` has reached the tail call, it is basically done. Any computations it needs to perform have been completed, and any local storage it has allocated (except possibly for variable `y`) is no longer needed. Instead of waiting until after `g` returns to deallocate its stack frame, it can do so before calling `g`. The compiler can then replace the call to `g` by a jump. Function `g` will then proceed in the normal manner, but when it executes its `ret` instruction, it will return directly to the function that called `f`.

Before jumping to `g`, the code for `f` must make sure that argument `y` is in the position that `g` expects. It can do this by storing the value of `y` at the address given by `%ebp + 8`. In doing so, it overwrites its own argument `x`, located in the caller's stack frame. Fortunately, this value is no longer needed.

```
1 /* Examples of function making sibling call */
2 int square(int x) {
3     return x * x;
4 }
5
6 int proc(int x, int y, int i) {
7     int b[2] = {x, y};
8     return square(b[i & 0x1]);
9 }
```

Figure 11: **Example of sibling call.** Function `proc` makes a sibling call to `square`.

Figure 11 shows an example of a function `proc` that makes a sibling call to a second function `square`. Normally, with optimization level two or higher, GCC will eliminate this call by inline substitution. By giving GCC the command-line argument '`-fno-inline`,' we disable inline substitution, and instead it will apply sibling-call optimization, yielding the following IA32 code for `proc`:

```
    Function proc, optimized -O3, inline substitution disabled
    x at %ebp+8
  1 proc:
  2   pushl   %ebp
  3   movl    %esp, %ebp
  4   subl    $16, %esp        Allocate 16 bytes on stack
    Create array b at %ebp-8 and set %eax to b[i & 0x1]
  5   movl    8(%ebp), %eax
  6   movl    16(%ebp), %edx
  7   movl    %eax, -8(%ebp)
  8   movl    12(%ebp), %eax
  9   andl    $1, %edx
 10   movl    %eax, -4(%ebp)
 11   movl    -8(%ebp,%edx,4), %eax
    Set up sibling call
 12   movl    %eax, 8(%ebp)     Set first argument to b[i & 0x1]
 13   leave                     Deallocate stack frame and restore frame pointer
 14   jmp     square            Make sibling call
```

This code follows the usual IA32 conventions, allocating 16 bytes on the stack for local array b and computing the argument for square in register %eax. Starting at line 12, we see the preparation for a sibling call. On this, line, the argument is stored at address %ebp + 8. On line 13, we see the current stack frame being deallocated with a leave instruction. On line 14, we see the jump to the code for square.

Sibling call optimization has only a minimal benefit in terms of the execution time of the resulting code. It replaces a call instruction by a jmp, and it eliminates one of the ret instructions, but these differences are minor. Unlike inline substitution, it does not enable a context-dependent optimization of the code for the sibling function. On the other hand, this optimization can yield a substantial savings in the total amount of stack space needed [1]. For example, consider functions f and g that call each other (mutual recursion) using tail calls. Normally, the stack space will grow in proportion to the depth of the recursion. With sibling-call optimization, the total stack space required remains constant, regardless of recursion depth.

7   Concluding Observations

Compiler-based code optimization is an active area of research, adapting to the changing performance characteristics of processors, as well as the ability to apply more sophisticated analyses and transformations by exploiting the increasing power of the processors on which the compilers are run. The developers of GCC continue to track this work, continually improving the ability of GCC to generate more efficient code.

Trying to debug a program compiled with sophisticated optimizations can be difficult, because the connection between the source code and the actual executable code becomes more tenuous. The best strategy is to keep the optimization level low (level one or none at all) until the program is ready for production usage.

Solutions to Problems

Problem 1 Solution: [Pg. 4]

A. For values of b between 9 and 15, it will insert $16 - b$ bytes. Otherwise, no bytes will be inserted.

B. For values of b between 1 and 7, it will insert $8 - b$ bytes. For values of b between 9 and 15, it will insert $16 - b$ bytes.

C. The combination of the two will have the same effect as the single directive `.p2align 3`. It is puzzling that GCC inserts both directives.

Problem 2 Solution: [Pg. 12]

Recursive factorial is another example of a linear recursion. We can map it to tail recursion in a manner similar to what we did for list summation.

A. Here is a tail-recursive version of factorial.

```
1 int fact_tail(int x, int sofar) {
2     if (x == 0)
3         return sofar;
4     else
5         return fact_tail(x - 1, sofar * x);
6 }
7
8 int fact_tail_call(int x) {
9     return fact_tail(x, 1);
10 }
```

The variable `sofar` accumulates the product of the successive values of `x`. The top-level function `fact_tail_call` initiates the computation with `sofar` equal to 1:

B. We can map the general form of tail recursion to this specific code as follows:

Abstract value	Program value
a_1	`x`
a_2	`sofar`
$Cond(a_1, a_2)$	`x == 0`
$Result(a_1, a_2)$	`sofar`
$Nval_1(a_1, a_2)$	`x - 1`
$Nval_2(a_1, a_2)$	`sofar * x`

C. Applying the transformation, we get the following C code:

```
1 int fact_itail(int x, int sofar) {
2     while (x != 0) {
3         int nx = x - 1;
4         sofar = sofar * x;
5         x = nx;
6     }
7     return sofar;
8 }
```

Problem 3 Solution: [Pg. 15]

The `iter1` transformation leads to a very conventional form of iterative factorial.

A. The following table shows the mapping between the general form and the factorial function:

Abstract value	Program value
a	x
a_0	0
b_0	1
$V(a)$	x
$N(a)$	x - 1
op	*

B. Applying the `iter1` transformation yields the following code:

```
1 int fact_iter1(int x) {
2     int r = 1;
3     while (x != 0) {
4         r = r * x;
5         x = x - 1;
6     }
7     return r;
8 }
```

Problem 4 Solution: [Pg. 16]

We can express the computation performed by this function by the following expression:

$$l_1 - [l_2 - [l_3 - [l_4 - \cdots [l_{n-1} - [l_n - 0)] \cdots]]]$$

A. We can rewrite $x - y$ as $x + -y$, and exploit the property that $-(-x) = x$, so that the expression can be written as the sum of alternating positive and negative terms. For the case where n is even, this yields the expression:

$$l_1 + -l_2 + l_3 + -l_4 + \cdots - l_{n-1} + l_n - 0$$

That is, the sum should alternate between adding the list element or its negation. A similar pattern holds when n is odd.

B. We can therefore rewrite the function, keeping its recursive form, but introducing a second argument
   wt that alternates between values $+1$ and -1. This is called at the top level with wt equal to $+1$:

```
1 int diff_list_rec_helper(list_ptr ls, int wt) {
2     if (ls == NULL)
3         return 0;
4     else
5         return (wt * ls->val) +
6             diff_list_rec_helper(ls->next, -wt);
7 }
8
9 int diff_list_rec2(list_ptr ls) {
10     return diff_list_rec_helper(ls, 1);
11 }
```

C. Applying a variant of the iter1 transformation yields the following iterative form:

```
1 int diff_list_iter1(list_ptr ls) {
2     int wt = 1;
3     int diff = 0;
4     while (ls != NULL) {
5         diff += wt * ls->val;
6         wt = -wt;
7         ls = ls->next;
8     }
9     return diff;
10 }
```

Problem 5 Solution: [Pg. 16]

This example illustrates the generality of the form given by function recur.

A. The key is to expand the computation V to include the call for the left subtree:

Abstract value	Program value
a	tp
a_0	NULL
b_0	0
$V(a)$	tp->val + sum_tree_rec(tp->left)
$N(a)$	tp->next
op	+

B. Applying the iter1 transformation gives the following code:

```
1 /* Sum values in tree,
2    recursively for left subtree and
3    iteratively for right subtree */
```

```
 4 int sum_tree_lrec(tree_ptr tp) {
 5      int sum = 0;
 6      while (!(tp == NULL)) {
 7          sum += tp->val +
 8              sum_tree_lrec(tp->left);
 9          tp = tp->right;
10      }
11      return sum;
12 }
```

GCC performs exactly this transformation when compiling sum_tree_rec.

Problem 6 Solution: [Pg. 17]

Here is the solution code. Since we are not assuming any special properties about op, it is important to combine values in the proper order.

```
recur_expand3(a) {
    r = b₀;
    if (a != a₀) {
        v[1] = V(a); /* First list element */
        a = N(a);
        if (a != a₀) {
            v[2] = V(a); /* Second list element */
            a = N(a);
            if (a != a₀) {
                v[3] = V(a); /* Third list element */
                a = N(a);
                r = recur_expand3(a);
                r = v[3] op r;
            }
            r = v[2] op r;
        }
        r = v[1] op r;
    }
    return r;
```

Problem 7 Solution: [Pg. 21]

This problem shows that commutativity is not a fundamental requirement for a transformation from recursion to iteration.

A. We can shift the parentheses to compute values from left to right, keeping the base value on the right:

$$[[\cdots[V(a) \; op \; V(N(a))] \; op \; \cdots \; op \; V(N^{k-2}(a))] \; op \; V(N^{k-1}(a))] \; op \; b_0$$

B. This leads to the following general form. We must treat the case where a equals a_0 explicitly:

```
iter2(a) {
    if (a == a₀)
        /* Special case when called with a₀ */
        return b₀;
    else {
        r = V(a);
        while (N(a) != a₀) {
            a = N(a);
            r = r op V(a);
        }
        /* Finish with a = a₀ */
        r = r op b₀;
        return r;
    }
}
```

Problem 8 Solution: [Pg. 21]

Applying the `iter2` transformation to recursive factorial gives this slightly unusual form:

```
1  int fact_iter2(int x) {
2      if (x == 0)
3          return 1;
4      else {
5          int r = x;
6          while (x-1 != 0) {
7              x = x - 1;
8              r = r * x;
9          }
10         /* Don't need to multiply by 1 */
11         return r;
12     }
13 }
```

Problem 9 Solution: [Pg. 21]

This is a very interesting and powerful transformation.

A. We know that if $N^k(a) = a_0$, then $P^k(a_0) = a$. We can therefore rewrite the expression as:

$$V(P^k(a_0)) \; op \; [V(P^{k-1}(b_0)) \; op \; \cdots \; op \; [V(P^2(a_0)) \; op \; [V(P(a_0)) \; op \; b_0]] \cdots]$$

B. This leads to a general form, where we start from a_0 and work up to a via P:

```
iter3(a) {
    r = b₀;
    if (a == a₀)
        /* Special case when called with a₀ */
        return r;
    else {
        t = P(a₀);
        while (t != a) {
            r = V(t) op r;
            t = P(t);
        }
        /* Finish with t = a */
        r = V(a) op r;
        return r;
    }
}
```

Problem 10 Solution: [Pg. 22]

This transformation leads to another common way to compute factorial iteratively.

A. Since subtraction is the inverse of addition, we can let $P(\texttt{x})$ be the operation x+1.

B. This leads to a version of iterative factorial where we count from 1 up to x:

```
1  int fact_iter3(int x) {
2      int r = 1;
3      if (x == 0)
4          return r;
5      else {
6          int t = 1;
7          while (t != x) {
8              r = t * r;
9              t = t + 1;
10         }
11         r = x * r;
12         return r;
13     }
14 }
```

References

[1] A. Bauer and M. Pizka. Tackling C++ tail calls. *Dr. Dobb's Digest*, February 1 2004.

[2] *GCC Online Documentation*. Available at `http://gcc.gnu.org/`.

[3] Y. A. Liu and S. D. Stoller. From recursion to iteration: What are the optimizations? *ACM SIGPLAN Notices*, 34(11):73–82, November 1999.

[4] S. S. Muchnick. *Advanced Compiler Design and Implementation*. Morgan-Kaufmann, 1997.

Index